Chapter 2: Replication

Lab - Writes with Failovers

Correct:

- When a `writeConcernError` occurs, the document is still written to the healthy nodes.

The `WriteResult` object simply tells us whether the `writeConcern` was successful or not - it will not undo successful writes from any of the nodes.

- The unhealthy node will have the inserted document when it is brought back online.

When the unhealthy node comes back online, it rejoins the replica set and its oplog is compared with the other nodes' oplogs. Any missing operations will be replayed on the newly healthy node.

Incorrect:

- `w: "majority"` would also cause this write operation to return with an error.

`w: "majority"` requests acknowledgement that a *majority* of nodes in a replica set have registered the write. In a three-node replica set, only two nodes are required for a majority, so the two healthy nodes are sufficient to satisfy this `writeConcern`.

- The write operation will always return with an error, even if `wtimeout` is not specified.

If `wtimeout` is not specified, the write operation will be retried for an indefinite amount of time until the `writeConcern` is successful. If the `writeConcern` is impossible, like in this example, it may never return anything to the client.

Proceed to next section