Final Exam

Question 4

Back to the Question

- **Indexes can solve the problem of slow queries.**

  This is correct.

- **Indexes are fast to search because they're ordered such that you can find target values with few comparisons.**

  This is correct.

- **Under heavy write load you should scale your read throughput by reading from secondaries.**

  No, since writes are replicated to secondaries all members of the replica set have about the same write workload, therefore sending reads to a secondary will not scale you read throughput.

- **When you index on a field that is an array it creates a partial index.**

  No, when you index a field that is an array it creates a **multikey** index.

- **On a sharded cluster, aggregation queries using $lookup will require a merge stage on a random shard.**

  No, `$lookup`, `$graphLookup`, `$facet`, and `$out` all require a merge stage on the **primary** shard, not a random shard like most other merged queries.

Proceed to next section